
Power to Choose? An Analysis of Consumer Inertia in the Residential Electricity Market

Author(s): Ali Hortagsu, Seyed Ali Madanizadeh and Steven L. Puller

Source: *American Economic Journal: Economic Policy*, November 2017, Vol. 9, No. 4 (November 2017), pp. 192-226

Published by: American Economic Association

Stable URL: <https://www.jstor.org/stable/26598351>

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <https://about.jstor.org/terms>



JSTOR

is collaborating with JSTOR to digitize, preserve and extend access to *American Economic Journal: Economic Policy*

Power to Choose? An Analysis of Consumer Inertia in the Residential Electricity Market[†]

By ALI HORTAÇSU, SEYED ALI MADANIZADEH, AND STEVEN L. PULLER*

Many jurisdictions around the world have deregulated utilities and opened retail markets to competition. However, inertial decision making can diminish consumer benefits of retail competition. Using household-level data from the Texas residential electricity market, we document evidence of consumer inertia. We estimate an econometric model of retail choice to measure two sources of inertia: search frictions/inattention and a brand advantage that consumers afford the incumbent. We find that households rarely search for alternative retailers, and when they do search, households attach a brand advantage to the incumbent. Counterfactual experiments show that low-cost information interventions can notably increase consumer surplus. (JEL D12, D83, L81, L94, L98, M31)

The deregulation of formerly regulated utilities has brought about more choice for energy consumers. Across the world, electricity and natural gas consumers increasingly have the ability to choose their retail provider. For example, households in over a dozen US states and over one-half of households in Europe have retail choice in electricity and/or natural gas. Consumers who previously purchased services from a utility are allowed to buy from other retailers at prices that are not regulated. This creates new markets where entrant retailers procure wholesale energy and market that energy to customers. This expansion of retail choice has been touted to have several benefits. Creating competition for the provision of utility services can lead to more competitive pricing in the short run. In addition, introducing competition can create incentives to provide customers with new value-added services.

However, choice frictions can diminish the benefits of retail choice. Households who never have had the power to choose may not exercise the option to select an alternative lower priced energy retailer. For example, households may not actively

*Hortaçsu: Department of Economics, University of Chicago, 1126 E. 59th Street, Chicago, IL 60637 (email: hortacsu@uchicago.edu); Madanizadeh: Sharif University of Technology, Azadi St., Azadi Square, Tehran, Iran (email: madanizadeh@gsme.sharif.edu); Puller: Department of Economics, Texas A&M University, TAMU 4228, College Station, TX 77843 (email: puller@econmail.tamu.edu). We thank the University of Chicago Energy Initiative, EI@Haas, NSF PSERC and NSF-SES 1426823 for generous financial support. We are grateful for data assistance from Kelly Brink, Robert Manning, Calvin Opheim, and Jess Totten. We thank Tim Brennan, Severin Borenstein, Koichiro Ito, Mar Reguant, Nancy Rose, Catherine Wolfram, three anonymous referees, and numerous seminar participants for very useful comments.

[†]Go to <https://doi.org/10.1257/pol.20150235> to visit the article page for additional materials and author disclosure statement(s) or to comment in the online discussion forum.

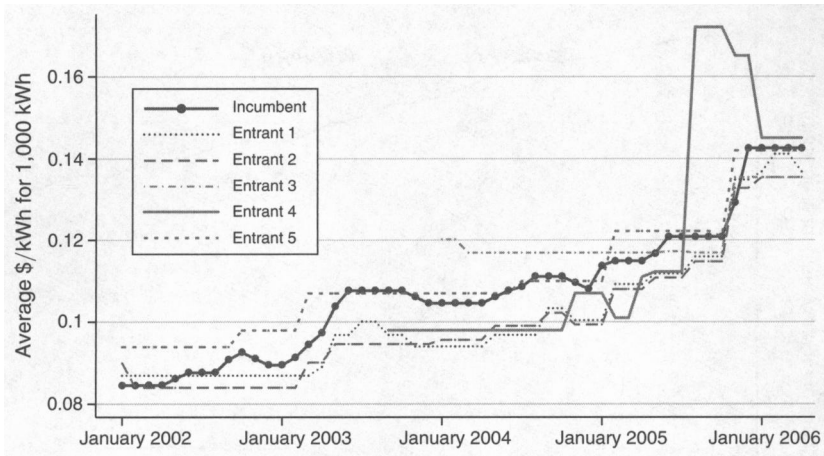


FIGURE 1. PRICES CHARGED BY INCUMBENT AND NEW ENTRANT RETAILERS IN FIRST FOUR YEARS OF MARKET

acquire information about other energy retailers, even if that information would indicate that better options exist. Alternatively, households may value the brand name of the incumbent—the old utility—and this may reduce the amount of switching to new entrant retailers. Both of these sources of choice frictions can reduce consumer gains of retail choice.

In this paper, we study a particular retail choice program to measure the size of choice frictions and to understand the underlying mechanisms. The Texas residential electricity market provides an excellent setting to investigate retail choice. Beginning in 2002, residential electricity customers were allowed to choose their retail provider. Initially, all households were by default assigned to the incumbent. In every subsequent month, households had the option to switch to one of several new entrant electricity retailers. In order to inform consumers and provide transparency to the search process, the Public Utility Commission of Texas created a website—*www.powertochoose.com*—that created one-stop shopping where households can search all retail options and switch to an alternative provider.

Aggregate data from this market suggest strong evidence of consumer inertia. Figure 1 plots the prices being charged by both the incumbent and new entrant retailers. Although prices varied over time, the incumbent's price was consistently higher than several of the new entrants. This suggests that households could reduce their electricity bills by switching from the incumbent to a new entrant retailer. However, Figure 2 shows that the incumbent held on to its market share leadership throughout the first four years of the market. The incumbent share only gradually eroded over time, and even after more than four years of retail choice, the incumbent's market share was over 60 percent. As we show in Section IIA, the majority of households did not switch away from the incumbent, even though they would have significantly reduced their electricity bills. Switching to a new entrant retailer—a one-time action that would take approximately 15 minutes to complete—would reduce the average electricity bill by around \$100 in the first year, which represents about 8 percent of electricity expenditures.




FIGURE 2. MARKET SHARES OF INCUMBENT AND NEW ENTRANT RETAILERS IN FIRST FOUR YEARS OF MARKET

This evidence of consumer inertia presents a puzzle for policymakers who are considering the deregulation of retail energy. Why do consumers of a relatively homogeneous product not switch to a lower priced energy provider? And can retail policy be designed to reduce consumer inertia? The goal of this paper is to understand the causes of the inertial behavior of residential electricity consumers. We develop a model to distinguish between two different mechanisms that could account for this inertia:

- *Search Frictions/Inattention*: Households may not choose to search for, or, equivalently, pay attention to offerings by other retailers. Even though www.powertochoose.com is a click away, the household may not be aware of it, or although informed of the existence of alternatives, the household chooses not to expend the effort to pay attention.
- *Incumbent Brand Advantage/Product Differentiation*: Even if a household is aware of other retailers, consumers may view the retailers, especially the incumbent, as vertically differentiated. For example, households may believe that service during power outage events or other dimensions of power quality could differ across retailers. Such beliefs, even if technically incorrect, may be a source of vertical product differentiation. Alternatively, consumers may believe that characteristics of customer support such as ease of paying bills vary across retailers.

Understanding the mechanism driving the inertial behavior will inform the design of policy to enhance the consumer benefits of retail choice. For example, if search frictions are the cause of observed inertia, then regulatory attempts to lower the cost of obtaining information about retail options could encourage switching and

benefit consumers. However, suppose that the primary reason households continue to purchase from the incumbent is that they believe the incumbent provides more reliable power. (Technically, the delivery of power and all metering operations are not a function of the retailer, but consumers may be unaware of this.) In this case, information campaigns to inform consumers that it's all the same power may induce households to choose lower price retailers. Finally, if the incumbent brand effect diminishes with time, then the resulting inertia may be viewed by policymakers as merely part of the consumer learning process during the transition from regulation to retail competition: while consumers may start with the belief that the incumbent delivers the superior product, over time they learn that it's all the same power.

In order to estimate the magnitude of each source of inertia, we develop an econometric model of household choice that nests both sources of inertia within a two-stage discrete choice framework. In each month, the household enters a two-stage process. In the first stage, the household decides whether to consider an alternative retailer. If it does not consider alternatives, then the household stays with its current retailer for the following month. However, if it does consider alternatives, the household enters the second stage. In this stage, the household observes the retailers in the choice set (which are available on www.powertochoose.com) and chooses the retailer that maximizes utility. In this second stage, we allow for vertical product differentiation to enter the household's decision, thus capturing potential brand advantage by the incumbent. In Section IIIC, we provide transparent conditions for the identification of model parameters from sample moments. We show that the first stage decision probabilities that capture inattention are separately identified from the second stage choice probabilities that capture incumbent brand advantage.

We estimate the model using household-level choice data from the first four years of retail choice in Texas. We find that both search frictions and the perception of brand differentiation explain the market dynamics shown in Figures 1 and 2. First, search frictions/inattention plays a role in the inertial behavior. We estimate that the average customer of the incumbent only searches for retail options in about 2 percent of months or approximately once every four to five years. However, the seasonality of search generates interesting insights about the determinants of consumer search in this market. Our model does not impose a structure on whether households are forward-looking when deciding to consider alternatives or, equivalently, the decision to not pay attention. Rather, households can either search in anticipation of seasonal patterns in consumption or as a reaction to a large bill caused by high consumption. We do not find evidence of forward-looking search activity: consumers are most likely to search in the month after receiving a large bill, which is likely to occur in the summer in Texas.

But inattention is not the only driver of inertia—the incumbent enjoys a very significant brand effect. In the earliest years of retail choice, consumers value the incumbent's brand at about \$60/month. However, this effect quickly diminishes over time and is estimated to be less than \$15/month by the fifth year of retail choice. These findings suggest a model of consumer learning where households who gather more experience with retail choice update their prior beliefs about the quality of the incumbent relative to new entrant retailers. Finally, we estimate our structural model separately for neighborhoods with different demographic characteristics. We

find that both sources of inertia are larger in neighborhoods with lower income, lower education levels, and more senior citizens. These findings suggest that retail choice policy can have important distributional implications.

In order to evaluate the implications for policy, we simulate the effect of a low-cost information intervention that targets each source of inertia. The information intervention serves to reduce inattention and mitigate misperceptions about the quality differences across retailers. This intervention would be akin to sending a flyer in the monthly bill telling consumers two important pieces of information: they can go to www.powertochoose.com to choose another retailer and that their power quality will be the same under any retailer. We view this as a fairly low-cost policy intervention that combines a nudge with an information treatment. This counterfactual is modeled as increasing the probability that a household searches in a given month and reducing the relative size of the perceived incumbent brand effect. We use model parameters to simulate this counterfactual experiment under various interpretations of the incumbent brand effect, and we find that consumer surplus could increase \$100 per year or more for each household.

A variety of countries have offered retail choice in utilities such as natural gas and electricity, and this has led to a mixed record on the number of consumers who switch away from the incumbent.¹ This paper contributes to the literature on studying retail choice in utilities. As one of the first jurisdictions to offer retail choice, the UK market has served as a valuable setting for analysis. Wilson and Waddams Price (2010) use rich survey data to study the demographics of switching and reported reasons for switching. They find that only a small fraction of switching customers choose the lowest priced retailer and that, in aggregate, switching consumers realize only between 30 and 52 percent of possible savings. Waddams Price, Webster, and Zhu (2013) use survey data to study demographic heterogeneity in search and switching behavior. In the most closely related paper, Giuliatti, Waterson, and Wildenbeest (2014) estimate an equilibrium model of search costs to explain price dispersion in the British electricity market. Retail choice has been investigated in natural gas (Giuliatti, Waddams Price, and Waterson 2005) and telecommunications (Miravete 2003). In studying choice behavior in these settings, Miravete and Palacios-Huerta (2014) have shown that it is important to consider the roles of both unobserved preference heterogeneity and state dependence in explaining choice persistence, as we discuss in Section IVC. Our study of deregulated retail markets expands upon a rich literature on deregulated electricity markets that has focused on wholesale markets (e.g., Wolfram 1999; Borenstein, Bushnell, and Wolak 2002; Bushnell, Mansur, and Saravia 2008; Sweeting 2007; and Hortaçsu and Puller 2008).

More generally, the phenomenon we study—frictions associated with allowing choice in settings where consumers previously did not have options—is not confined to formerly regulated utilities. In the health care sector, the prescription drug benefit program under Medicare Part D provides the elderly with multiple private plan options rather than a single plan specified by the government; likewise, health

¹For an analysis of the merits of retail choice in electricity, see Joskow (2000). For a review of experiences in residential electricity choice and a representation of consumer search in such markets, see Brennan (2007).

exchanges are an alternative to employer-provided insurance.² In primary education, parents in some jurisdictions are offered a menu of public schools that their children can attend rather than a single school that children are zoned to attend.³ And in the arena of retirement savings, the traditional role of government in pay-as-you-go systems has been replaced by privatized retirement planning where individuals choose from among a set of privately managed funds.⁴

Our paper is also related to a recent literature developing methods to estimate preferences in settings where decision making is influenced by search frictions/inattention. Hong and Shum (2006), Hortaçsu and Syverson (2004), and Moraga-González and Wildenbeest (2008) are early attempts that utilize aggregate market level data, and more recent efforts by, e.g., Kim, Albuquerque, and Bronnenberg (2011); De los Santos, Hortaçsu, and Wildenbeest (2012); De los Santos, Hortaçsu, and Wildenbeest (2013); Honka (2014); Koulayev (2014); and Honka, Hortaçsu, and Vitorino (2014) utilize detailed consumer level data on both choices and the search process/consideration sets (as obtained from website clicks) to test and estimate models of consumer search. Our empirical setting is one where we observe the choices of consumers but do not observe their search process/consideration sets. However, we show that the search frictions/inattention component of consumer behavior and the “frictionless” preference component can be separately identified under reasonable assumptions. Indeed, as we show in our results section, applying standard discrete choice models to our data without taking into account the presence of search frictions/inattention can yield distinctly different and implausible estimates of preferences. Several papers in the recent empirical literature on consumer inertia exploit the institutional feature that some customers are new to the market while others are existing market participants (e.g., Handel 2013 and Luco 2014). In some settings, new customers face different sources of inertia than existing customers, and comparing the decisions of these different groups allows for the identification of inertia. In other institutional settings (including our setting), there may be relatively few new customers to exploit for identification. As we describe below, our empirical strategy does not require new customers, although it can incorporate new customers if they are present.

I. Retail Electricity Choice in Texas

Residential electricity customers in Texas, as in many states, historically were served by vertically integrated utilities at regulated prices. The state was divided into separate service territories, each with a vertically integrated firm. Beginning in 2002, residential electricity customers in Texas were allowed to choose their retail provider. In January 2002, all customers by default were assigned to a retail firm that was affiliated with the old incumbent utility. In any subsequent month, a customer can switch to any other retailer at no cost. Texas is not alone in this

²For example, see Abaluck and Gruber (2011); Kling et al. (2012); Ketcham et al. (2012); and Dafny, Ho, and Varela (2013).

³See, for example, Hastings, Kane, and Staiger (2010); Neilson (2013); and Hastings and Weinstein (2008) on school choice and the effects of information provision.

⁴See Hastings, Hortaçsu, and Syverson (2013); Duarte and Hastings (2012); and Luco (2014).

regulatory change—over a dozen US states have opened the retail electricity sector to competition.

Any customer choosing to procure power from another retailer is entering into a *financial* agreement with the retailer. Importantly, retail choice does not impact the *technical* operations of power provision. The former utility was split into a “lines” company and a retailer.⁵ The operations of all electricity transmission, local powerlines, and meters is now operated by a regulated firm (with a different name) that is a separate business entity from the incumbent retail firm. As a result, the quality of power service (e.g., outages) is independent of the retailer chosen by a household. (It is possible that consumers were not aware of the independence of operations, as we discuss below when interpreting our results).

The incumbent’s price was regulated by the public utility commission and called the “price-to-beat.” The incumbent could request an adjustment to this regulated rate up to two times a year; however, the size of the adjustment was indexed to natural gas input costs.

The new entrant retailers are firms that procure power in the wholesale market and sell retail power to residential customers. These retailers need not own any physical infrastructure to be market participants. Unlike the price of the incumbent, the prices of the entrant retailers were not regulated. In 2002, most parts of Texas had between three to five entrant retailers, and by the end of our sample, in 2006, the choice set expanded beyond ten.

Because of relatively low wholesale electricity prices during the first few years of retail competition, the price-to-beat was considered to be higher than competitive prices for retail power. This so-called “headroom” was an intentional market design feature by the regulatory commission to ensure new retailers of sufficient price-cost margins to encourage entry. As a result, entrant retailers were able to price more than one cent per kilowatt-hour (kWh) below the price-to-beat. As we discuss below, this created the potential for average savings of about \$12 per month for switching away from the incumbent.

Households had multiple sources of information on potential electric retailers. The most salient source of information was a well-publicized website established by the state’s public utility commission—www.powertochoose.com. A screenshot of the website is shown in Figure II in the online Appendix. Consumers enter their zip code and view a list of all retailers. Households could follow links on this website to switch to a new retailer in a process that was online and relatively quick. A majority of consumer search appears to have occurred through this website, as we document with hitcount data in Section IVB.⁶

The public utility commission sought to reduce the costs of switching in several ways. First, the incumbent was not permitted to charge the customer to switch to another retailer.⁷ In addition, a household switching away from the incumbent could

⁵ The generation (i.e., powerplant) sector had already been separated from the utility prior to 2002.

⁶ In contrast, information acquisition in the UK market appears to have been through face-to-face channels with only 10 percent of customers using price comparison websites (Wilson and Waddams Price 2010).

⁷ The incumbent was required to offer a month-to-month plan, so customers faced no dollar switching costs and no lock-in. However, the entrant retailers were allowed to offer longer term contracts with early termination fees that were required to be waived if the customer moved. As we describe below, our data does not include information

switch back to the incumbent's price-to-beat later if it chose to do so. Regardless of the retailer chosen, a household would still receive a single monthly bill that included charges for all electricity services (energy, transmission, distribution, metering, and billing). Finally, any customer moving to a different residence after the beginning of retail choice is required to choose a retailer; the default is no power. Less than 1 percent of observations in our sample are movers.

II. Data

We study the retail choice behavior of all residential electricity customers in the service territories covered by one of the formerly vertically integrated regulated utilities—Texas New Mexico Power (TNMP). This utility has the appealing feature that it was formed by several mergers and therefore has customers spread throughout the state, including both urban and rural areas.

We use monthly data on each of the approximately 192,000 residential meters in TNMP territory from January 2002 until April 2006. For each meter, we have information on the electric retailer used by the household and the electricity consumption for every month. We use the meter address to match to census block group data on demographics.

We focus on customers of the six retailers that had at least 1 percent of the market at some point in time from January 2002–April 2006.⁸ These six retailers are the incumbent, two retailers that were new in the TNMP service territory but were affiliated with incumbents from other parts of the state, and three retailers that were new to the Texas marketplace.

We can calculate each household's bill with high precision. For each household-month, we have data on the household's retailer and its electricity consumption. We match these data to information on the tariffs that were charged by each retailer in a given month.⁹ All of the tariffs for the new entrant retailers were 12-month fixed rates; early termination fees were waived by law if the customer moved. The incumbent's tariff was the regulated price-to-beat which would only change if the incumbent requested an adjustment based on natural gas prices. The tariff data allow us to calculate the total electric bill that a household paid to its chosen retailer as well as counterfactually how much the household would have paid to purchase the same amount of power from any of the other retailers.¹⁰

at the customer level on contract duration. So our estimates of the behavior of households served by entrant retailers may reflect, in part, any lock-in during the first year of service. However, few customers switch more than once during our sample period, suggesting that the effect of lock-in with new entrants is unlikely to be large.

⁸We exclude one entrant retailer with 1.1 percent market share for which no price data are available.

⁹The Public Utility Commission of Texas collected monthly information on the rate plans offered by each retailer in the different service territories. In many cases, the retailer offered only a single rate plan, so we can precisely measure the monthly bill. In fact, only one rate plan was offered by four of the six retailers that we model, including the incumbent. However, a complication is presented by the fact that two retailers offer a menu of rate plans, and we have no information on which plan is chosen by a given household. In these two cases, we chose the plan that was thought to be most popular by the analyst at the public utility commission with responsibilities of overseeing the retail market.

¹⁰We should note that about 6 percent of customers received discounts as part of a low-income program. However, eligibility for this program was independent of retailer. We do not have data on which customers qualified for this program, so we are forced to assume that these customers pay the standard tariff.

Most of the retailers use multipart, or nonlinear, tariffs.¹¹ Because the rate plan is an important determinant of consumer choices, we need to choose an appropriate measure of price that is likely to drive a household's decision process. Two options are available—the marginal price and the average price. The marginal price is likely to be the same for all households with any given retailer because even those retailers with increasing block tariffs have the highest block begin at a low usage level (400 kWh/month). However, the average price differs by (expected) consumption because all but one retailer employ nonlinear tariffs.

Although some research on nonlinear tariffs assumes that consumers respond to the marginal price, the assumption is problematic in this setting. Responding to the marginal price requires households to have full information on the multipart tariff function. This assumption is not likely to hold in this market. First, *powertochoose.com* saliently displays only the *average* price for customers consuming 1,000 kWh/month. Figure II in the online Appendix shows a sample screenshot. Consumers who want to gain more detailed information may click to download the "Facts Label" that is required to contain specific parameters of a retailer's service. The rate information on the Facts Label is the average price for customers consuming either 500, 1,000, or 1,500 kWh; however, no information was available on the shape of the nonlinear tariff at other consumption levels. Thus, consumers did not have the ability to calculate the cost of purchasing their specific usage from any of the retailers.

Recent empirical work suggests a second reason that marginal price may not be a suitable metric of the price that affects consumer choice. Ito (2014) studies residential electric customers in California who face different nonlinear tariffs and finds evidence that customers respond more to the average than to the marginal price. More generally, utility bills in the United States typically display information that make the average price (total bill due divided by usage) much more salient than the underlying nonlinear tariff schedules.

For this reason, when we develop a discrete choice model in Section III, we use the retailer's *average price at 1,000 kWh/month* as the metric of price for several reasons. First, the *powertochoose.com* website saliently displayed the average price at 1,000 kWh as the price used to sort retailers. Second, this usage level is close to the typical usage level of customers in our sample—the average consumption is 1,140 and the median consumption is 968 kWh/month. Third, we favor using the average price over the customer's expected bill size because information to calculate bill size for the customer's usage was not available.

Figure 1 shows the evolution of the average price at 1,000 kWh by each retailer. Rates ranged from about 8.5 to 14 cents/kWh from 2002–2006. Rates were generally rising with much of this driven by increases in the price of natural gas, a primary determinant of wholesale electricity prices in Texas. In particular, rates

¹¹For the incumbent, the tariff is a fixed fee of \$5.17 and then an increasing block tariff with the second block beginning at 400 kWh per month. Three of the other retailers had similar tariff structures—fixed fees around \$5 and increasing block tariffs beginning in the middle of the sample period. One retailer had a relatively higher fixed fee of \$8.70 followed by a single block tariff. And the last retailer had a linear tariff—no fixed fee and a single block tariff.

jumped several cents in late 2005 following the natural gas price spikes caused by Hurricanes Katrina and Rita.

The average price of the incumbent (i.e., the price-to-beat rate) was systematically higher than one or more entrants throughout most of the sample period. In fact, by the fall of 2002, the first year of retail choice, at least one entrant offered an average rate at least one-half cent cheaper than the incumbent in every month except one month in late 2005. Moreover, in many months in the middle of the sample, an entrant's average rate was over one cent cheaper than the incumbent's price-to-beat.

Some households move between residences during our sample. In order to identify movers, we use information on disconnect/reconnect status for each meter.¹² The number of household-months involving a move comprises less than 1 percent of all observations in our sample. We discuss how we include movers into the analysis in Section IVC.

A. Summary Statistics

First, we provide a basic description of observed switching behavior and the dollar magnitude of expenditure reductions of purchasing from alternative retailers. These summary statistics yield patterns that are consistent with results that arise from our model's estimates in Section IV, thus providing support for our modeling assumptions.

Switching: Timing and Frequency.—The incumbent maintained a large market share despite charging higher rates than entrant retailers. Figure 2 shows the market shares of the six largest retailers over the first four years of retail choice. The incumbent exhibited a slow erosion of market share throughout the sample, but still maintained over a 60 percent share by April 2006. Two other retailers had over a 10 percent share while the remaining retailers had less than a 5 percent share.

The total number of switches per month was relatively low in the first year of retail choice but then rose in the following three years. Figure 3 displays the total number of switches from one retailer to another for each month. There appears to be a seasonality in switching behavior with a peak in the summer months when electric bills can be twice as high as winter bills. Our model below allows for seasonality in the decision to consider alternative retailers so that we can explore this channel of decision making.

Finally, we measure the frequency with which a household switches retailers. Approximately two-thirds of households never switch away from the incumbent. Among those households that do switch, most switch only once (21 percent of meters) or twice (9 percent of meters). A histogram of the number of times that a given household switches retailers during the first four years of retail choice is shown in the online Appendix in Figure III.

¹²Specifically, our data on the monthly choice of retailer is reported for the meter/address. We do not know the names of the customers on the bills, so we only directly measure when the residence is served by a different retailer. To identify movers, we assume that the resident at the meter does not change unless there is a disconnect of service. If electric service is disconnected for more than 30 days, we assume that new residents occupy the residence. Otherwise, we assume the same residents are making decisions for the residence.

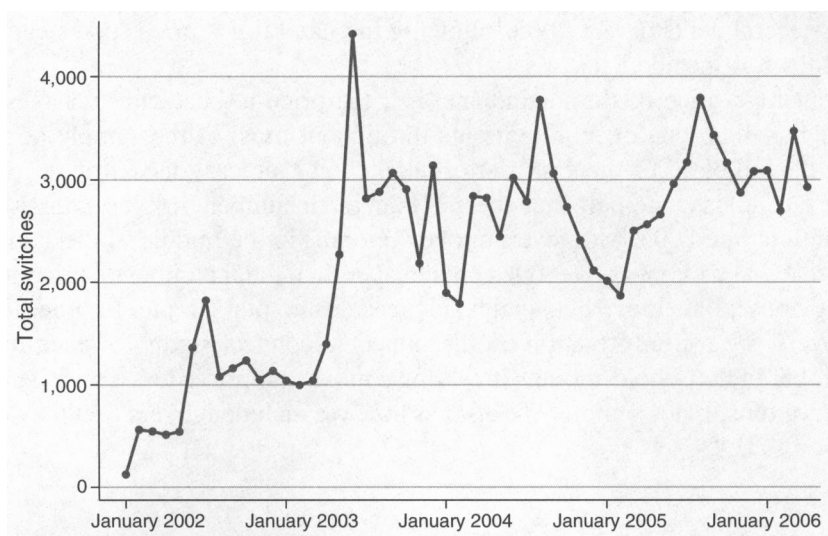


FIGURE 3. TOTAL NUMBER OF SWITCHES OF RETAILER BY MONTH

Potential Savings from Switching.—Either of the two potential sources of inertia—search frictions/inattention or an incumbent brand effect—could cause consumers to purchase from the incumbent while lower priced retail options are available. A descriptive metric of inertia is the dollar savings to consumers of purchasing from an entrant retailer rather than the incumbent. This section measures the potential dollar savings from switching. These savings should be viewed as *descriptive* evidence; we develop a model to decompose the sources of inertia in Section III.

We measure the savings to households of buying the same amount of power from an alternative retailer. To do so, we calculate the bills for each household in months it purchased from the incumbent and the counterfactual bill if the household had purchased the same amount of power from other retailers. As noted above, these reductions in expenditures should be not seen as consumer surplus calculations.

Consider two extremes of the frequency with which a household switches. First, consider a scenario in which a consumer switches *only once* during the four year sample and does so in the first month of retail choice (January 2002). We calculate the monthly savings if each household had switched to one of the two large entrant retailers. The mean savings of purchasing from one of the larger retailers is \$7.69/month and the mean savings purchasing from the other is \$9.97/month.

At the other extreme, consider a scenario in which a consumer switches to the lowest price retailer *each* month. For households buying from the incumbent, the mean savings per month of switching to the lowest price retailer is \$12.47, and the median savings is \$7.63.¹³

¹³ We also assess the savings realized by those households that switched rather than purchased from the incumbent, using similar assumptions as above. For those months in which households purchased from any retailer other than the incumbent, the average bill would have been \$8.86 higher per month if the same consumption were purchased from the incumbent.

These savings suggest that there is sizable inertia to switching in the early years of this market. Households could purchase the same amount of the same power for less—averaging between roughly \$7–\$12 dollars per month, or \$84–\$144 dollars if scaled up to one year. This savings represents about 8 percent of *total* electricity expenditures. Our model below estimates the extent to which consumers continue to purchase from the higher priced incumbent because consumers do not consider alternative retailers or because the consumers view the incumbent as a higher valued product.

These metrics could mask considerable heterogeneity across different demographic subpopulations. In the online Appendix, we break down savings from switching by demographic group. We find that more of the potential savings of switching is realized by households in neighborhoods that are wealthier, have higher education levels, and fewer senior citizens.

Factors Associated with Switching Away from Incumbent.—Finally, we document the relationship between switching behavior and key variables in the model we present below. Specifically, we show that a household is more likely to switch away from the incumbent when recent monthly electricity bills are large and when there are more entrant retailers with lower prices than the incumbent. To show these correlations, we estimate a linear probability model using data for each household-month for the sample period used in our structural model below. We focus on household-months when the household was served in the previous month by the incumbent, and we estimate the probability of switching away from the incumbent. The baseline probability that a household served by the incumbent switches to another retailer in a given month is 1.0 percent.

Table 1 reports correlations based on regressions that include fixed effects for each household. Thus, we exploit within-household variation in the probability of leaving the incumbent for an entrant retailer. Column 1 shows that a household is more likely to switch away from the incumbent when there are more entrant retailers with lower prices than the incumbent. The presence of one additional entrant with lower prices increases the probability of switching by 15 percent relative to the baseline switch rate. In column 2, we show that households are more likely to switch after receiving a large monthly bill. Differences in a customer's bill are primarily driven by the large seasonality in electricity consumption. A doubling of the last bill increases the switch rate by 0.27 percent or about one-quarter of the baseline rate. Column 3 shows that there is seasonality in switching, even after controlling for the number of lower priced retailers and the size of the most recent bill. Column 3 includes calendar quarter fixed effects, and we find that switching is more likely *during* and *after* the summer demand peak as compared to the lead-up to the peak in the spring.

These relationships show that switching is associated with both past information received by the household—the size of the last bill—and by characteristics of the plans offered by entrants. Next, we use these relationships to develop a model of switching.

III. Model

In this section, we build an econometric model that allows us to separately identify the two sources of inertia—search frictions/inattention and an incumbent brand

TABLE 1—DESCRIPTIVE ANALYSIS OF SWITCHING AWAY FROM INCUMBENT

	Dependent variable: Indicator of switching from incumbent		
	(1)	(2)	(3)
Number of cheaper entrant retailers	0.0015 (0.0001)	0.0016 (0.0001)	0.0020 (0.0001)
log of last monthly bill received		0.0027 (0.0001)	0.0010 (0.0001)
Calendar quarter 2			0.0038 (0.0001)
Calendar quarter 3			0.0069 (0.0001)
Calendar quarter 4			0.0070 (0.0001)
Constant	0.0047 (0.0002)	−0.0081 (0.0005)	−0.0056 (0.0005)
Household fixed effects	Yes	Yes	Yes
Observations	3,729,919	3,729,919	3,729,919

Notes: This table reports factors that are associated with switching away from the incumbent. An observation is a household-month when the household was served by the incumbent in the previous month. The dependent variable is an indicator of whether the household switched away from the incumbent to an entrant retailer in that month; the mean switch rate is 1 percent. We estimate the correlations with a linear probability model using household fixed effects.

effect. We model the household-level choice of electricity retailer as a two-stage process that occurs each month. Each month, in stage one, the household has a current retailer and decides whether to consider alternative retailers with some probability. We refer to this stage as the decision-to-choose stage. In stage two, the households who choose not to consider alternative retailers in stage one will maintain their current retailer for the following month.¹⁴ However, households who choose to consider alternative retailers in stage one will then choose the retailer that maximizes utility among those in the market. These households may choose a different retailer or may continue with their current one. We refer to this second stage as the choice stage. We allow for heterogeneity across households and across time at both the decision-to-choose stage and the choice stage, as we describe below.

One empirical complication is that we do not observe the outcome of the decision-to-choose stage. We only observe households who change retailers, i.e., those who decide to consider alternative retailers *and* then choose a different one. From the analyst's viewpoint, households who do not switch are both those who do not consider alternatives and those who do consider alternatives but choose their current retailer. We describe a model below that allows us to separately identify parameters of the decision stage and the choice stage.

A. Stage One: Decision to Choose

Each month a household decides whether to consider alternative retailers. We model the probability of considering a (possibly new) retailer to vary by

¹⁴ Again, we do not take a stance as to whether the choice to consider alternatives is rational or forward-looking.

characteristics of the household. First, the probability varies by the household's current retailer. For example, a household's experience with its existing retailer may induce it to consider alternatives. Anecdotal evidence from industry analysts suggests that households are driven to consider alternative retailers in response to events such as a large summer bill or by a bad experience with the current retailer. This envisions the decision to choose as a push rather than a pull effect.

Second, we allow for seasonality in the months of the year that customers actively decide upon their retailer. This allows the model to attribute some of the rise in switching during summer months observed in Figure 3 to result from more searching during the summer. In some specifications, we allow the decision probability to vary in household characteristics.

We denote the *current* retailer by k and the *new* retailers by j (again, recall that k and j can be the same if the household searches and chooses its current retailer). The month of the sample is indexed by t .

We model the decision probability for any household that is currently a customer of retailer k at time t , denoted λ_t^k , with a standard binary logit:

$$(1) \quad \lambda_t^k(\gamma) = \frac{e^{W_t^k}}{1 + e^{W_t^k}},$$

where $W_t^k = \sum_r \gamma_r Z_{rt}^k$, and $\{Z_{rt}^k\}$ is a set of observable characteristics including dummy variables for each existing retailer k and month-of-year dummy variables to allow for seasonality in deciding to search. In some additional specifications, we also include the dollar change in the size of customer i 's most recent bill as compared to the previous bill in order to capture an increase in salience due to a large bill. (In those specifications, the λ_t^k and W_t^k have i subscripts, but we do not include them here).¹⁵

We view this specification of the decision to choose as a reduced-form representation of the drivers of inattention. This is a critical feature of modeling choice in settings with inertia due to inattention. As we show in Section IVC, if one were to exclude this stage of the model and only use a standard discrete choice model in which households chose from the choice set *every* month, then one would make incorrect inferences about preference parameters.

Our descriptive evidence in Section II suggests that modeling the impediment to search with a model of *inattention* is appropriate in this setting. In other settings, researchers have used *switching costs* to explain choice frictions. Switching costs can be viewed as generating choice frictions—via state dependence—even if buyers are fully informed about other options. In contrast, search cost/inattention represents factors such as psychic search costs and not having information, or being willing to gather information, about alternatives. There are several reasons that switching costs are unlikely to be primary drivers of inertia in this market. First, there are no monetary costs to switching and powertochoose.com makes time costs small. That being said, it is possible that consumers were unaware of the rule that prohibited switching

¹⁵This formulation can be seen as a reduced-form representation for an underlying search protocol or rational inattention model.

fees and the presence of the easy-to-use website. Therefore, customers may have had expectations of meaningful switching costs, even if such costs do not exist, and those expectations deter search. The data are not consistent with full information about options—one of the key features of a switching cost model. To see this, recall that the data strongly suggest seasonal patterns in both search and switching. (As we document in Section IIA, switching peaks in the summer, and as we discuss later in Section IVB, the number of visitors to www.powertochoose.com also peaks in summer). It is difficult to explain why a large bill would induce search if households had perfect information about alternatives.

B. Stage Two: Choice of Retailer

In this stage, each household who enters the choice stage in month t chooses the retailer from the choice set that yields the highest utility, as in a standard discrete choice model. For households who enter the choice stage, the indirect utility for each household i of choosing retailer j in period t is

$$(2) \quad U_{ijt} = V_{ijt}(\theta) + \varepsilon_{ijt},$$

where $V_{ijt}(\theta)$ is a parameterized utility term, and ε_{ijt} is a random utility shock that is i.i.d. across consumers, retailers, and time. We assume ε_{ijt} to be a Type I Extreme Value random variable. The term $V_{ijt}(\theta)$ is further specified as $V_{ijt}(\theta) = \sum_s \theta_s X_{jt,s}$ where $X_{jt,s}$ is the s th characteristic of retailer j at time t . The product characteristics include the price, an indicator for whether the retailer is the incumbent, and the incumbent indicator interacted with a linear time trend. This specification allows for vertical product differentiation—a brand effect—by the incumbent. Specifically, the variables comprising X_{jt} are:

- p_{jt} is the price of retailer j in month t for 1,000 kWh usage per month (as reported on the website powertochoose.com). As we discuss in Section II, the average price is arguably the most salient metric of price that affects choice. Note that by using the current price, we assume that consumers expect future retail prices to reflect current prices. This assumption is consistent with Anderson, Kellogg, and Saltee (2013) who find that consumers' beliefs about another energy commodity—gasoline—are consistent with a no-change forecast. This allows us to view household-level choice as a static model.
- *INCUMBENT* is an indicator variable for the incumbent retailer, allowing for an incumbent brand effect.¹⁶
- *INCUMBENT* · *MONTHCOUNTER* is the incumbent indicator interacted with the number of months since the market began, allowing for a linear time trend in the incumbent brand effect.

¹⁶In principle, we could estimate differential brand effects across the entrant firms, but we do not observe enough switching among entrants to do so. Product differentiation among entrants is thus captured by the logit shocks.

Because ε_{ijt} is a Type I Extreme Value random variable, the probability that household i chooses retailer j in month t is given by the familiar logit probability:

$$(3) \quad P_{ijt}(\theta) = \frac{\exp(V_{ijt}(\theta))}{\sum_{k \geq 1} \exp(V_{ikt}(\theta))}.$$

This probability is used in GMM estimation that is described below.

C. Simultaneously Estimating Decision and Choice Stages

We simultaneously estimate both the decision to consider alternative retailers (stage one) and choice (stage two). In order to do so, we exploit our data on observed switching behavior to derive a set of moment conditions. As noted above, one empirical challenge is that we do not directly observe the outcome of stage one. Rather, we observe switches to other retailers for those who decide to consider alternative retailers. Thus, households who do not switch can be either households who did not consider alternatives or households who searched and chose to stay with their existing retailer.

To address this complication, we exploit the observed month-to-month aggregate switching from the old retailer k to the new retailer j to estimate the probability of search. This model of the flow of customers from one retailer to another provides moments for our GMM estimation. First, we provide a simple example to illustrate the empirical strategy and then we present the formal model.

Illustration of Empirical Strategy.—We illustrate the empirical strategy with a simple example that also allows us to show how we can separately identify the search and decision stages. Assume that we only observe two months of data—the customer's retailer last month and this month. In addition, assume that there are only three retailers. Let each household served the previous month by retailers 1, 2, and 3 decide to consider alternative retailers with probability given by λ_1 , λ_2 , and λ_3 , respectively. And let the probability of choosing retailer 1, 2, and 3 conditional upon entering the choice stage, be given by P_1 , P_2 , and $1 - P_1 - P_2$, respectively. We want to estimate these five probabilities, $\{\lambda_1, \lambda_2, \lambda_3, P_1, P_2\}$.

Conceptually, we can create a matrix of counts of the number of customers switching from retailer k to retailer j during the month. See Figure 4 for an illustration. The cells of this matrix provide us with statistical moments that we use to estimate the five parameters/probabilities.

Consider all households who were served by retailer 1 the previous month, and denote this number $N^{(1)}$. Some of these households will be observed to use retailer 1 in the current month; such households are ones who did not consider alternative retailers (occurring with probability $1 - \lambda_1$) and those who considered alternatives but chose to remain with retailer 1 (occurring with probability $\lambda_1 \times P_1$). Likewise, consider households observed to use retailer 1 in the previous month and retailer 2 in the current month. These households are ones previously with retailer 1 who considered alternatives and chose retailer 2 (occurring with probability $\lambda_1 \times P_2$). Likewise, we can characterize households previously with retailer 1 who considered alternatives and chose retailer 3 (occurring with probability $\lambda_1 \times (1 - P_1 - P_2)$).




FIGURE 4. ILLUSTRATION OF IDENTIFICATION STRATEGY

The *expected* number of customers who were initially with retailer 1 and continue to use retailer 1 is: $N^{(1)} \cdot [(1 - \lambda_1) + \lambda_1 P_1]$. The *expected* number of customers who were initially with retailer 1 and switched to one of the other retailers is: $N^{(1)} \cdot \lambda_1 P_j$ for $j = 2, 3$. This provides 3 moments to match to sample moments on the number of customers flowing between retailer 1 the previous month and the 3 retailers in the current month.

One can derive similar sets of 3 moments for customers who were initially served by retailer 2 and retailer 3. This provides us with 9 moment conditions. However, one moment in each set is redundant because any customer who does not stay with, say retailer 1, must switch to retailer 2 or 3. Thus, we exclude the redundant moments for customers who stayed with the same retailer and use the off-diagonal terms of the flow matrix as the moments for estimation. This yields six moments to estimate the five probabilities of interest.

This simple example provides the basic intuition for estimating the model. In our setting, we have more than two months of data so we have $T - 1$ flow matrices and could estimate decision and choice probabilities for each period. In practice, we parameterize the decision and choice probabilities to specific variables of interest, as described in Sections IIIA and IIIB.

Formal Specification of Empirical Strategy.—We formalize the estimation strategy described above and allow for households to be heterogeneous in their search and choice probabilities. Let $B_t^{(k)}$ be the set of households whose retailer was k at time $t - 1$. Let $N_t^{(k)}$ denote the total number of households in $B_t^{(k)}$. Further, λ_{it}^k is the probability that a household in $B_t^{(k)}$ chooses to search in period t . For those households who choose to search, P_{jt} is the probability that the household chooses j in period t . Finally, at the end of the two-stage process each month, $N_{jt}^{(k)}$ is the total number of households in $B_t^{(k)}$ who use retailer j at time t .

Suppose household i is in the set $B_t^{(k)}$. Let $d_{ijt}^{(k)}$ be an indicator function of whether household i (who had been served by retailer k in time $t - 1$) is served

by retailer j at time t . The expected value of this indicator variable prior to period t is:

- For $j \neq k$ (households changing retailers):

$$(4) \quad E_{t-1}[d_{ijt}^{(k)}] = \frac{\lambda_{it}^k P_{jt}}{\Pr(\text{deciding and choosing } j)}$$

- For $j = k$ (households keeping the same retailer):

$$(5) \quad E_{t-1}[d_{ikt}^{(k)}] = \frac{\lambda_{it}^k P_{kt}}{\Pr(\text{deciding and staying with } k)} + \frac{1 - \lambda_{it}^k}{\Pr(\text{not deciding})}$$

Because $N_{jt}^{(k)}$ is the total number of households in $B_t^{(k)}$ with retailer j at time t , we have

$$(6) \quad N_{jt}^{(k)} = \sum_{i \in B_t^{(k)}} d_{ijt}^{(k)}.$$

Thus, our moment equations tell us that the expected number of customers previously served by k who are now served by j is given by

- For $j \neq k$:

$$(7) \quad E_{t-1}[N_{jt}^{(k)}] = \sum_{i \in B_t^{(k)}} \lambda_{it}^k P_{jt}.$$

- For $j = k$:

$$(8) \quad E_{t-1}[N_{kt}^{(k)}] = \sum_{i \in B_t^{(k)}} (\lambda_{it}^k P_{kt} + 1 - \lambda_{it}^k).$$

The last flow equation (8) showing the flow from k to k is redundant because the probabilities of moving away from k and staying with k add up to 1. Thus, we use the off-diagonal ($j \neq k$) moments for estimation. This yields $J(J-1)$ moments for each time t .

We use GMM to estimate (γ, θ) , the parameters determining decision (λ_{it}^k) and choice ($P_{jt}^{(k)}$) using the objective function:

$$\min_{\gamma, \theta} \eta' W \eta,$$

where $\eta \equiv \langle \eta_{jt}^{(k)} \rangle$ and $\eta_{jt}^{(k)} = \frac{N_{jt}^{(k)} - (\sum_{i \in B_t^{(k)}} \lambda_{it}^k P_{jt})}{N_{jt}^{(k)}}$, and W is a weighting matrix.

The intuition behind the objective function is straightforward. Consider the numerator of each moment. The term $N_{jt}^{(k)}$ is the number of households in our data that switch from retailer k to retailer j in month t . Our model says that the expected number of households switching from k to j is $\sum_{i \in B_t^{(k)}} \lambda_{it}^k P_{jt}$. We plug in the specified functional forms for the decision and choice probabilities, given in

equations (1) and (3). GMM finds parameters that make the model most closely fit the data on the number of switchers. The denominator simply downweights moments with larger variance; it adjusts for the fact that the number of customers for which the model is off is likely larger for retailers that have a large number of customers.

Identification.—The identification argument is a generalization of the simple example considered above and illustrated using Figure 4. The matrix capturing the flow of customers from retailer k to retailer j allows for separate identification of the probabilities of search λ^k and the probabilities of choice P_j . Mathematically, this matrix provides moments of the order J^2 while we are only estimating probabilities on the order J . Embedded in this model are two key assumptions. First, the decision probability is a function of the last retailer k but not the next retailer j . As we note above, this envisions a push rather than a pull model of search, and is consistent with views of industry analysts. Second, the choice probability is a function of the next retailer j and not the last retailer k . This assumption implies that upon deciding to consider alternative retailers (e.g., on www.powertochoose.com), consumers consider all retailers on equal terms and do not have private information on any retailers based on their past experience. One might be concerned that prior experience with a retailer influences choice (for reasons beyond observed product characteristics). However, keep in mind that very few customers switch multiple times during our sample period, so the vast majority of customers have experience with only one or at most two retailers. Finally, it is important to note that our identification strategy allows for the possibility that large bills during the summer induce consumers to search for alternative retailers but not necessarily switch unless the entrant prices are sufficiently lower than the incumbent price.

Importantly, this identification strategy does not require us to have choice data for both consumers with preexisting market experience and consumers new to the market. Several papers in the existing literature on inertia exploit the fact that some customers are new while others have preexisting market experience (e.g., see Handel 2013 and Luco 2014). The presence of new market participants—movers in our case—fits into our modeling approach, but we do not require the movers in order to separately identify inattention bias from vertical product differentiation. To see this, a mover in our setting must choose a retailer (otherwise they have no power), so the probability of considering retailers in stage one must equal one ($\lambda_{it} = 1$). Non-movers may or may not consider alternative retailers in any given month ($\lambda_{it} \leq 1$). Any customer who considers alternative retailers—whether a mover or non-mover—then enters into the stage two choice of retailer.

Therefore, our model can be estimated where movers' decisions are given by $\lambda_{it} = 1$ and the choice probability given by equation (3) and where non-movers' decisions are given by λ_{it} as specified in equation (1) and the choice probability from equation (3). But it is also clear that the parameters of the model are identified even if no movers are present (i.e., we could use only non-movers—those existing customers who can switch from retailer k to retailer j —to estimate the model parameters using equations (1) and (3)). In our setting, we have only a small number of

movers—less than 1 percent of our observations—so an empirical strategy that does not depend on new market participants is valuable.

IV. Results

We organize our results into several sections. First, we show estimates of our benchmark model that estimates how often incumbent customers search and how much those customers differentially value the incumbent brand when they do search. Then, we allow the probability of searching in any given month to vary so that we are able to identify important patterns in inattention bias. Following the main estimates, we illustrate important features of our modeling framework in Section IVC. Next, in Section IVD we estimate the model separately for neighborhoods with different demographic characteristics in order to identify heterogeneities in how consumers respond to retail choice. Finally, these results are used to motivate our counterfactual policy experiments that we present in Section V.

For all of the results presented below, we estimate the model on a 20 percent random sample of meters in our data in order to ease the computational burden. In addition, we restrict attention to the period of January 2004–April 2006 when all six firms that we analyze are present in the market. In most of our results, we include all customers who currently have a retailer and study their decision making to choose another retailer; however, we show results that include new customers (movers) in Section IVC.

A. Benchmark Estimates of Inertia

In our benchmark specification, we model the stage one decision of household i in period t to consider alternative retailers to be a function of whether the household's existing retailer is the incumbent or not. Then conditional upon deciding to search, the stage two choice model is a logit model where each product is characterized by the price of consuming 1,000 kWh. We include an incumbent brand effect and allow the brand effect to have a linear time trend.

Our benchmark model results are reported in column 1 of Table 2.¹⁷ The table reports the parameters of both the stage one decision step and the stage two choice step. In addition, we use the parameter estimates to calculate other metrics that aid in the economic interpretation of these parameters. Specifically, for the decision step, we calculate the probability that a customer with a given firm chooses to consider alternative retailers in a given month. Also, we use the choice parameters to calculate price elasticities of each firm (evaluated at the average price) and the dollar value of the incumbent brand effect.

The parameters of the decision-to-choose model indicate that households do not frequently search. Specifically, when the parameters of λ_{it}^k are converted to probabilities, our model estimates that customers of the incumbent only search in 1.8 percent of months. This identifies a large source of inertia in the market. For many households, who by the design of retail choice were defaulted to the incumbent, a search will

¹⁷In online Appendix Table II, we report counts of the number of switchers that serve to identify our parameters.

TABLE 2—PRIMARY MODEL RESULTS

	Benchmark (1)	Seasonality in search (2)	Large bill affects search (3)	All (4)
<i>Stage one: Decision to choose</i>				
Parameters (γ)				
Constant	−3.363 (0.045)	−3.720 (0.123)	−3.468 (0.048)	−3.643 (0.149)
Incumbent	−0.643 (0.064)	−0.647 (0.062)	−0.589 (0.068)	−0.626 (0.064)
January		0.222		0.038
February		0.375		0.315
March		0.266		0.043
April		0.041		−0.125
May		0.210		0.176
June		0.228		0.170
July		0.638		0.364
August		0.635		0.474
September		0.541		0.424
October		0.547		0.429
November		0.383		0.365
Large bill (\$ change in two most recent bills)			0.007 (0.002)	0.003 (0.003)
Estimated effects				
Pr(monthly search) incumbent customer (λ)	0.018	0.018	0.018	0.017
Pr(monthly search) entrant customer (λ)	0.033	0.033	0.032	0.032
<i>Stage two: Choice of retailer</i>				
Parameters (θ)				
Price (cents/kWh)	−0.435 (0.091)	−0.464 (0.083)	−0.462 (0.096)	−0.445 (0.085)
Incumbent brand dummy	2.764 (0.256)	2.946 (0.269)	2.789 (0.294)	2.834 (0.293)
Incumbent month-of-sample counter	−0.076 (0.014)	−0.086 (0.015)	−0.075 (0.016)	−0.080 (0.016)
Estimated effects				
Incumbent price elasticity	−2.52	−2.61	−2.67	−2.55
Average entrant price elasticity	−4.51	−4.82	−4.80	−4.62
Incumbent brand effect (\$/mo) in January 2004	\$61.86	\$61.61	\$58.72	\$61.85
Incumbent brand effect (\$/mo) in April 2006	\$14.87	\$11.67	\$14.66	\$13.50

Notes: This table reports results from the benchmark structural model of Section III estimated via GMM. Parameter estimates are reported with standard errors in parentheses for the parameters of the two stages of the model. Then the table reports point estimates of model effects that are calculated using the parameter estimates, in order to facilitate model interpretation. We do not report standard errors for the month fixed effects in the interest of space.

not occur until many months into the new market. This probability can be used to calculate that only 19 percent of households have searched at least once within one year of the market opening, 35 percent within two years, and 61 percent by the end of our sample over four years after retail choice begins.¹⁸ Thus, our model implies that inattention bias is an important driver of inertia. In our counterfactual experiments in Section V, we estimate the impact of policies that increase search.

Our model also estimates the probability that a customer of an entrant retailer considers alternative retailers. We find that the search rate for customers of entrants—

¹⁸This calculation assumes a simple i.i.d. structure to the decision-to-choose process; in later specifications, we allow for a richer process.

searching in 3.3 percent of months—is larger but still relatively small. One possible reason that these customers may search more is a selection effect; these customers necessarily have searched at least once before if they are customers of an entrant retailer.

The parameters of the choice model show that the incumbent brand effect is another source of inertia. The positive coefficient on the incumbent brand dummy variable indicates that, conditional upon deciding to choose, customers attach higher utility to the incumbent's product than to the entrant retailers. Notably this brand effect declines with time, as indicated by the negative coefficient on the brand effect interacted with a month-of-sample counter. We monetize the size of this brand effect using the coefficient on the price. As shown in the last rows of the table, the brand effect is \$61.86 per month in January 2004.¹⁹ This means that after accounting for price differences between the incumbent and entrants, consumers value purchasing from the incumbent nearly \$62 *per month* more than purchasing the same power from any of the entrant retailers. This brand effect is identified by the extent to which decreases in the price differential between the incumbent and entrants leads more consumers to choose an entrant retailer. This initial incumbent brand effect is quite large, corresponding in dollar terms to nearly half on a customer's typical monthly bill. However, it is important to keep in mind that the size of the brand effect is not bounded by expenditures. This product differentiation yields differences in the own-price elasticities of demand. As shown in the table, the incumbent's price elasticity is -2.52 while the entrant retailers have elasticities averaging -4.51 .

Importantly, the size of this brand effect declines as more months pass since the beginning of retail choice. This declining brand effect is identified by the rate of change of searching customers who choose the incumbent versus new entrants as the price differential between the incumbent and entrants narrows with time (see Figure 1).²⁰ By the end of our sample, in April 2006, the incumbent brand effect for customers still with the incumbent is significantly smaller; it has declined to \$14.87 per month. Given the type of selection at work, the consumers who remain with the incumbent after more than two years of retail choice are those with higher brand effects. We find that even for these customers, the brand effect declines substantially with time. This suggests that the large incumbent brand effect does not persist after the market matures and customers have more experience with retail choice.²¹

Several plausible explanations exist for this sizable brand effect. One possibility suggested to us by market analysts is that customers believe that the quality of their power depends on the firm providing electric service. As we discuss above, this is not the case—retail choice is a financial relationship between customer and retailer that

¹⁹This uses the standard logit model approach to calculating consumer surplus. The price is measured in cents/kWh. We use the incumbent brand coefficient, the time trend in that coefficient, and the price coefficient to estimate in dollars the brand effect for 1,000 kWh of usage equals $((2.764 + (-0.076) \times t)/0.435) \times (1,000/100)$ for $t = 1$ (January 2004) to 28 (April 2006).

²⁰Consider the two entrants with the largest market shares—entrant 1 and 2 in Figure 1. These two entrants' average price was \$9.66/month cheaper than the incumbent in January 2004 (for 1,000 kWh). By the beginning of 2006, the two entrants' average price was only \$6.40/month cheaper.

²¹We also estimate a version of our model where both the incumbent brand effect and incumbent search rate are allowed to vary linearly with time. In this specification, the incumbent brand effect in January 2004 is \$61.80 and in April 2006 is \$14.46, yielding a similar decline in the incumbent brand effect as in our benchmark model. The time trend in search by incumbent customers is always small but actually falls over time.

has no impact on the physical distribution of power. Because the physical transmission and distribution system is operated by a firm that is independent from retailers, a customer will not see any difference in factors such as power outages, restoration of power in the event of an outage, or meter reading services. However, customers may not have been aware of this fact because of an incomplete understanding of the market.

Consumer learning is a related interpretation of the very large incumbent brand effect at the beginning of deregulation followed by its gradual erosion. Consumers may have started out with the belief that the incumbent is the more reliable provider. However, over time, either through observational learning or by social interactions, they may have revised their priors via Bayesian learning toward believing that service quality is not differentiated. It is beyond the scope of this paper to model the evolution of the brand effect; rather, we focus on implications of a brand effect for policy that allows retail choice.²²

Another possible source of the incumbent brand effect may have been a fear that entrant retailers would go bankrupt and leave them without service. Technically, if a retailer exits the market, customers are automatically shifted back to the incumbent at a price set at 130 percent of the wholesale price; however, this may have been unknown to customers. Finally, customers may have feared bait-and-switch tactics by new entrant retailers. Although such tactics are possible, we do not see such an expectation borne in the data.²³

Unfortunately, our data does not allow us to isolate the exact mechanism driving the estimated brand effect and its evolution over time. However, in our counterfactual experiments in Section V, we estimate the effect on consumer surplus if the relative size of the brand effect (whatever the cause) were made smaller.

B. When Do Households Search?

The descriptive analysis in Section IIA shows that there are strong seasonal patterns in both bill size and switching behavior, with peaks occurring during the summer. In our next specifications, we allow for temporal and consumer level heterogeneity in the probability to decide to search so that we can better understand potential determinants of search. Specifically, we quantify the seasonality of search and the extent to which a customer receiving a large bill induces search. Results are shown in columns 2–4 of Table 2.

The first new specification allows for seasonality in the decision to search by adding separate dummy variables for each month of the year to the decision probability λ_{it}^k . As seen in column 2, the coefficient of the monthly dummy variables are substantially higher during and immediately after the summer months, implying that search behavior is most intense at the end of the summer.

We can validate this result of the model with outside data on search. Data on the number of visitors to www.powertochoose.com were provided to us by the Public Utility Commission for the period 2005–2009. (Unfortunately, data covering our

²² A rich literature in marketing has explored identification in models of Bayesian learning including Erdem and Keane (1996) and Shin, Misra, and Horsky (2012).

²³ As discussed above, the rise in prices late in our sample period was driven by wholesale cost shocks.

sample period were not recorded). In these later years, the pattern of website visitors exhibited similar seasonal patterns, with the highest number of visits occurring in August. Thus, our model findings, corroborated with outside evidence, suggests that large electric bills may make electric choice more salient and induce households to assess their retail options.²⁴

In order to explore the effect of bill size, we estimate another specification that incorporates whether the household recently received a large bill. We modify our benchmark model so that the decision probability λ_{it}^k includes a household-specific measure of the dollar difference between the most recently received monthly bill and the previous bill. This variable—Large Bill—primarily reflects increases in consumption rather than changes in prices. As shown in column 3, Large Bill has a positive and statistically significant affect on the decision probability, suggesting that receiving an unusually large electric bill induces customers to search for alternative retailers. In column 4, we include both Large Bill and monthly dummies simultaneously. It appears that much but not all of the Large Bill effect is driven by seasonality of electricity consumption.

This relationship between recent bill size and the probability of searching offers insights into the determinants of household search behavior. A priori one might expect that households would search for lower priced retailers in the spring *in anticipation* of the summer peak in electric consumption and expenditures. The potential savings from switching to a lower priced retailer is highest if the switch occurs before the summer. However, we do not find evidence of anticipatory search. Rather, consumers appear to react to large increases in summer consumption and respond by searching. Overall, we do not find evidence that consumers are forward-looking when choosing the time to search.

The results in columns 2–4 also suggest that our estimates of inertia are robust to different specifications of the decision to search probability. The estimates of the two major sources of inertia—the average decision probability and the incumbent brand effect—are quantitatively very similar to the estimates from the benchmark model.

C. Illustrating the Implications of Our Modeling Framework

In this section, we compare our two-stage model of choice to other empirical methodologies that study choice in the presence of inertia. First, we show that in our setting, our main conclusions are not affected by the empirical challenge of separating consumer heterogeneity from state dependence. Second, we show that if one were to use a standard discrete choice model—the second stage of our two-stage model—then one would obtain results that are not plausible, illustrating the value of our two-stage model.

²⁴We also use the hitcount data to validate anecdotal evidence that a majority of search was conducted via the website. Based on the search rates estimated by our model and the fraction of customers statewide served by incumbents and new entrant retailers, we estimate that approximately 112,000 customers each month considered alternative retailers (i.e., entered stage two of the model). The monthly average number of unique visitors to powertochoose.com was 77,000 customers. While keeping in mind that these comparisons are approximations, we believe they provide strong suggestive evidence that the powertochoose.com website was the primary channel of search behavior during our sample period.

Incorporating Movers to Address Heterogeneity versus State Dependence.—As we discuss in the introduction, some papers in the existing literature have exploited new market participants that face different levels of inertia. Our model does not require new customers for identification, but it can incorporate them into the model.

Having data on movers is useful in assessing whether the “heterogeneity versus state dependence” problem is an important issue in our setting. As has been documented earlier, one reason consumers appear inertial may be due to selection based on unobserved preferences that are persistent over time (e.g., Keane 1997; Dube, Hitsch, and Rossi 2010; and Miravete and Palacios-Huerta 2014). Thus, not accounting for selection based on persistent unobservable preferences may bias the decomposition of observed temporal choice patterns into a preference component versus “inertia” component.

The advantage of our institutional setting, similar to the setting in Handel (2013), is that some households are movers and these households do not have the initial condition problem when choosing a retailer. The reason is that a large majority of these movers are coming from other states or other parts of Texas where the incumbent utility is another firm. (Eighty percent of movers to counties in the service territory we study came from areas where the incumbent is another firm, based upon IRS data documenting the origin and destination counties of migrating tax filers.)

These movers allow us to assess the possible role of state dependence in affecting our estimates of the incumbent brand effect. If there is state dependence in the primary sample that we use (i.e., the non-movers), this could bias upwards our estimate of the incumbent brand effect. Thus, if there is a problem due to state dependence, we would expect our estimates of the brand effect *using only movers* to be smaller.

We utilize our data on movers in two ways. First, we incorporate moment conditions for the subsample of movers into our estimation. Specifically, we model new customers—or movers in our setting—as not facing inattention because they must choose a retailer to get power to their house, but they still may face an incumbent brand effect. We include observations in which a household transitioned from no retailer to some retailer, generating new moments for GMM given by equation (7). The probability of entering the choice stage is set equal to one ($\lambda_{it}^{mover} = 1$). The probability of choosing retailer j is given by equation (3).

We allow the preference parameters to differ between movers and non-movers because the incumbent brand effect could be quite different for the two groups. On one hand, some customers moving from outside the service territory of the incumbent may not know the incumbent and thus have a smaller incumbent brand effect. On the other hand, some movers may come from locations that do not have retail choice, which could make those customers more attached to the incumbent per se, regardless of the identity of that incumbent. Results are shown in Table 3. Column 1 redisplayes the results of our benchmark model (the same as in column 1 of Table 2). Column 2 shows estimates from the model that includes both non-movers and movers. For the non-movers, our estimates of the two sources of inertia are unchanged. For the movers, we estimate an initial incumbent brand effect that is about 20 percent higher than for non-movers. However, this brand effect declines substantially slower for movers—the brand effect is \$75/month in January 2004 and only falls to \$52/month by April 2006. Thus, while movers do not face inertia

TABLE 3—ILLUSTRATING IMPLICATIONS OF OUR MODELING FRAMEWORK

	Benchmark (1)	Include movers (2)	No decision stage (3)
<i>Stage one: Decision to choose</i>			
Parameters (γ)			$\lambda_{it}^k == 1$
Constant	−3.363 (0.045)	−3.363 (0.045)	
Incumbent	−0.643 (0.064)	−0.643 (0.064)	
Estimated effects			
Pr(monthly search) incumbent customer (λ)	0.018	0.018	$== 1$
Pr(monthly search) entrant customer (λ)	0.033	0.033	$== 1$
<i>Stage two: Choice of retailer</i>			
Parameters (θ)			
Price (cents/kWh)	−0.435 (0.091)	−0.435 (0.091)	−0.008 (0.002)
Incumbent brand dummy	2.764 (0.256)	2.764 (0.256)	0.132 (0.024)
Incumbent \times month-of-sample counter	−0.076 (0.014)	−0.076 (0.014)	−0.004 (0.001)
Mover \times price		−0.090 (0.118)	
Mover \times incumbent		1.231 (0.276)	
Mover \times incumbent \times month-of-sample counter		0.030 (0.015)	
Estimated effects			
Incumbent price elasticity	−2.52	−2.51	−0.08
Average entrant price elasticity	−4.51	−4.46	−0.08
Incumbent brand effect (\$/mo) in January 2004	\$61.86	\$61.86	\$163.51
Incumbent brand effect (\$/mo) in April 2006	\$14.87	\$14.87	\$30.80
Movers incumbent brand effect (\$/mo) in January 2004		\$75.25	
Movers incumbent brand effect (\$/mo) in April 2006		\$51.58	

Notes: This table reports results to illustrate the implications of our modeling framework, as described in Section IVC. Column 1 reports our benchmark model results, which is also column 1 from Table 2. Column 2 includes new market participants (movers) by restricting movers to choose in stage one ($\lambda_{it}^k = 1$) and allowing movers to have different preference parameters, as described in Section IVC. Column 3 illustrates the need for a decision stage of the model by removing the decision stage and imposing that $\lambda_{it}^k = 1$.

Source: Authors’ calculations

from inattention, they continue to favor the higher priced incumbent due to placing additional value on purchasing power from the firm known as the incumbent retailer in their area.²⁵

To confirm the estimated brand effect for movers, we go a step further and reestimate our model using only movers in order to directly address the “heterogeneity versus state dependence” decomposition. We model a mover as entering the choice

²⁵In the online Appendix, we show supplementary evidence that this result is not model-driven but rather can be seen in raw data. Specifically, we document the market share of movers in the first month after a move when the customer must choose a retailer. The incumbent share is very high, despite the incumbent charging higher rates, but the incumbent share declines over time.

stage with probability equal to one ($\lambda_{it}^{mover} = 1$) and then choosing a retailer by the stage two discrete choice logit model. As shown in the online Appendix in Table III, there is still a sizable brand effect—the estimated brand effect at the beginning of the period is estimated to be \$79 for the movers (this compares to our estimate of \$62 for non-movers). Importantly, for these moving households who do not have an initial conditions problem, the brand effect is larger than our estimated brand effect for the non-movers, which is opposite what one would expect if there were consumer heterogeneity induced by state dependence.

A second piece of evidence suggesting this is not a concern in our data is the time trend in the estimated brand effect. If the persistence of purchasing from the incumbent were due to state dependence, then the households remaining with the incumbent over time are likely to have higher average $\varepsilon_{i,INCUMBENT}$. Under these circumstances, our model that does not account for these correlated shocks would yield a brand effect that *increases* over time. In contrast, we find an estimated brand effect that *falls* over time. Thus, while in general one must be concerned with the heterogeneity versus state dependence problem noted in Keane (1997) and addressed by Miravete and Palacios-Huerta (2014), the above results suggest that the problem is not going to affect our main conclusions in our particular setting.

Illustrating the Need for a Decision Stage.—Next, we explore the implications of modeling retail switching with our two-stage process and why a standard one-stage discrete choice framework would yield results that are implausible. Our decision to model switching as a two-stage process is motivated by the belief that consumers in this market do not actively choose retailers every month. Thus if we were to estimate a standard one-stage discrete choice model, we believe that we would make incorrect inferences about consumer choice.

To illustrate this point, we apply a standard discrete choice model to our data. We operationalize this by imposing that the decision probability $\lambda_{it}^k = 1$; this is effectively turning off the stage one decision step and assuming that all households enter the stage two choice step. Under this assumption, in every month where we observe a customer to stay with her current retailer, that customer is finding that the current retailer's product maximizes utility. Column 3 of Table 3 shows estimates where we turn off the stage one decision step.

Comparing these two sets of results, we see two major differences. First, the implied incumbent brand effect under no decision step is substantially larger (e.g., \$164 as compared to \$62 in January 2004). This is not surprising because imposing that consumers search every period means that the model must rationalize staying with the higher priced incumbent entirely with a brand effect.

Second, and more importantly, the own-price elasticities of residual demand for the entrant retailers are implausibly small (around -0.08). Such a low own-price elasticity contradicts standard oligopoly theory in which profit-maximizing firms price in the elastic portion of their residual demand. (Recall that the incumbent could not choose its price, but the new entrants were free to set any price.) These estimated price elasticities strongly suggest model mis-specification. Both of these results illustrate the need to choose an empirical model that allows for consumers to infrequently search for alternative retailers.

D. Does Inertia Vary across Demographic Groups?

The size of consumer inertia in markets with retail choice can vary in important ways across different parts of the population. As shown in the online Appendix, the amount of realized dollar savings from switching varies notably across neighborhoods with different demographic characteristics. The descriptive analysis shows that the fraction of savings realized is nearly twice as large in neighborhoods that were more wealthy, more educated, and have few senior citizens. In this section, we test if the two sources of inertia vary across neighborhoods.

We estimate our benchmark model separately for different subpopulations based upon income, education, and age. We classify a household based upon whether its census block group is above or below the median census block group when ordered by household income, fraction of the population with at least a bachelor's degree, and the fraction of the population that is over age 65. Results are reported in Table 4.

We find that both sources of inertia are larger for neighborhoods with lower income, lower education, and more senior citizens. Specifically, customers of the incumbent consider alternatives at different rates, as estimated by the parameter $\lambda_{\text{INCUMBENT}}$ in Table 4. Customers of the incumbent search for alternative retailers with higher frequency in neighborhoods with a higher income by 2.3 versus 1.3 percent (columns 1–2), in neighborhoods with a more college educated population by 2.6 versus 1.0 percent (columns 3–4), and in neighborhoods with fewer senior citizens by 2.4 versus 1.3 percent (columns 5–6).²⁶

We also find that the *initial* incumbent brand effect is larger in certain neighborhoods; however, the effect declines to very similar levels by the end of our sample. In neighborhoods above the median in income, the brand effect is \$52/month in January 2004 while it is \$113/month in below median income neighborhoods. However, this vertical differentiation declines at rates such that both types of neighborhoods have similar brand effects by 2006 (\$13 and \$23). When splitting neighborhoods by education, the below median education neighborhoods have a very large brand effect at the beginning of the sample (\$267/month), but this brand advantage declines to \$17/month by 2006. In contrast, the higher educated neighborhoods have an initial brand effect of only \$47/month, which declines to \$20. Finally, neighborhoods with more senior citizens have an initial brand effect that is larger than the brand effect in younger neighborhoods, but the difference disappears by 2006. These results suggest that retail choice policy can have important distributional consequences, particularly in the early years of choice. We discuss policy implications after presenting our policy counterfactuals in Section V.

This convergence in the size of the incumbent brand effect across different types of customers provides suggestive evidence of consumer learning. Even though households with different demographics begin with different priors on the brand value of the incumbent, households converge on a very similar value after several

²⁶ These demographic differences coincide with heterogeneity in customer engagement found in the UK market. The Competition and Markets Authority (2015) reports, based on detailed survey data, that households less likely to switch retailers, less likely to consider switching, and less likely to have shopped are households with older residents, lower income, and less education.

TABLE 4—BENCHMARK MODEL BY NEIGHBORHOOD DEMOGRAPHICS

	Income		Education		% senior citizens	
	Below median	Above median	Below median	Above median	Below median	Above median
<i>Stage one: Decision to choose</i>						
Parameters (γ)						
Constant	−3.298 (0.100)	−3.385 (0.048)	−3.300 (0.104)	−3.437 (0.039)	−3.426 (0.047)	−3.292 (0.093)
Incumbent	−1.019 (0.151)	−0.350 (0.066)	−1.284 (0.155)	−0.197 (0.055)	−0.291 (0.071)	−1.021 (0.148)
Estimated effects						
Pr(search) if incumbent customer (λ)	0.013	0.023	0.010	0.026	0.024	0.013
Pr(search) if new retailer customer (λ)	0.036	0.033	0.036	0.031	0.031	0.036
<i>Stage two: Choice of retailer</i>						
Parameters (θ)						
Price (cents/kWh)	−0.275 (0.146)	−0.499 (0.104)	−0.134 (0.111)	−0.456 (0.095)	−0.497 (0.093)	−0.406 (0.235)
Incumbent	3.204 (0.538)	2.641 (0.275)	3.707 (0.602)	2.194 (0.173)	2.327 (0.225)	3.616 (0.507)
Incumbent \times month-of-sample counter	−0.091 (0.030)	−0.071 (0.015)	−0.124 (0.032)	−0.046 (0.009)	−0.055 (0.012)	−0.104 (0.028)
Estimated effects						
Incumbent price elasticity	−1.35	−3.04	−0.59	−2.98	−3.23	−1.81
Avg. entrant price elasticity	−2.93	−5.13	−1.46	−4.67	−5.08	−4.36
Incumbent brand effect (\$/mo) in January 2004	\$113.31	\$51.51	\$266.78	\$47.10	\$45.70	\$86.41
Incumbent brand effect (\$/mo) in April 2006	\$23.47	\$13.13	\$17.08	\$19.89	\$16.04	\$17.05

Notes: This table reports results of estimating the benchmark model (column 1 of Table 2) split by the demographic characteristics of the household’s census block group. A household is classified by whether its census block group is above or below the median among all census block groups. Education is defined by the fraction of the population with a bachelor’s degree or above.

Source: Authors’ calculations

years. This is consistent with a simple model of Bayesian learning in which consumers have different initial beliefs about the relative quality of the incumbent. However, the differences in the priors are eventually washed away as data accumulates from observational learning or social interactions.

V. Policy Counterfactuals:
The Effects of an Information Intervention to Reduce Inertia

Our model estimates the size of two sources of inertia. The mechanisms that inhibit switching away from the higher priced incumbent are that the incumbent’s customers only search in 2 percent of months, and when they do search, customers place a sizable brand effect on the incumbent’s service. Next, we estimate how much consumer surplus would increase under an information intervention that reduces the size of each source of inertia.

Our hypothetical policy intervention is targeted to households who are inertial—those who continue to purchase from the incumbent after two years of retail choice

in January 2004. Our intervention has two dimensions. The first dimension is to inform inertial customers that they have the ability to choose their retailer and tell them where they can go to find a list of retailers and each retailer's offering. The second dimension is to inform customers that their power quality is entirely independent of their retailer. Specifically, purchasing from another retailer is buying power that is equally reliable from a technical point of view. As we discuss above, the brand effect captured by our model could include other dimensions of quality such as customer service, but there is strong anecdotal evidence that much of this brand effect captures the perception that the incumbent provides more reliable power.

Practically, this intervention could be an informational flyer. Suppose that the regulator required the incumbent to attach a one-page flyer with the January 2004 monthly bill that prominently displayed two pieces of information:

- "The State of Texas has created a website www.powertochoose.com where you can see all the options available to you. It's quick. It's easy to use. And you can switch your retailer at no cost to you in 15 minutes or less."
- "*It's all the same power*—the quality of electrical service will not change because <<Firm X>> controls your powerlines rather than <<The Incumbent>> or any other retailer."

We view this information intervention having two effects on customer decision making. First, it will increase the probability of a customer searching for alternative retailers. Customers who stay with the incumbent for two years of retail choice may exhibit inattention, as described in Section IIIA. Also, customers may expect that they will face time or dollar costs of switching to another retailer, even if in reality those costs are negligible. Informing customers about the easy-to-use website and the fact that there is no cost to switch is likely to adjust expectations about switching costs and encourage search. Second, it will reduce the relative brand advantage of the incumbent. Of course, this information intervention is only a nudge—households still choose the retailer that maximizes utility.

The magnitude of the effect of this information intervention is an empirical question that we cannot directly assess without actually conducting a randomized controlled trial. Therefore, we conduct counterfactual calculations under different assumptions to provide a range of estimates of the consumer surplus effects of this low-cost information intervention.

The information intervention is modeled as reducing inertia via two factors: increasing the number of incumbent customers who search and changing consumer perceptions about the size of the incumbent brand effect. The size of the consumer surplus effect depends in part on the interpretation of the brand effect. On one hand, consumers may attach direct utility to being a customer of a specific firm. On the other hand, the brand effect may be a factor that distorts choices but does not directly generate utility. For example, if the brand effect arises from the incorrect perception that power quality differs across firms, then one could view this effect as a factor that should not enter a social planner's welfare calculation.

We do not have information that allows us to isolate the primary driver of the estimated brand effect. Therefore, we estimate lower and upper bounds of changes

in consumer surplus under different interpretations of the brand effect. The lower bound models the incumbent brand effect as distorting choice behavior but not affecting ex post utility. In our upper bound calculation, the brand effect changes both choices and the utility of buying from a given retailer. As we show below, even the lower bound estimate suggests that this low-cost information intervention can meaningfully increase consumer surplus.

In our upper bound, the information intervention affects the perceived value of the entrant retailers relative to the incumbent. We calculate the expected change in consumer surplus of moving from the status quo to the counterfactual information intervention. The change in expected consumer surplus per household is given by $\Delta E(CS)_i = \frac{1}{-\alpha} \left[\ln \left(\sum_{j=1}^{J^{CTF}} e^{V_{ij}^{CTF}} \right) - \ln \left(\sum_{j=1}^{J^{SQ}} e^{V_{ij}^{SQ}} \right) \right]$ (Small and Rosen 1981), where *CTF* denotes counterfactual and *SQ* denotes status quo. This is the difference in log-sum terms divided by the price parameter that converts utils to dollars.²⁷ First, the intervention is modeled as increasing the probability that a household searches for alternative retailers in a given month. Under the status quo, 2 percent of the incumbent's customers consider among *all* retailers and choose the retailer that maximizes utility; 98 percent of customers choose from a set that only includes the incumbent. Under the counterfactual with less inattention, we increase the fraction of incumbent customers who consider all $J = 6$ alternatives and reduce the fraction who only "choose" the incumbent from the $J = 1$ choice set. Second, the intervention is modeled to reduce the relative brand advantage of the incumbent. We conceptualize the information treatment that "It's all the same power" to increase the brand effect of the entrant retailers to some fraction of the brand effect of the incumbent. To calculate this effect, we augment the product characteristics of each entrant retailer to have its own brand effect that is some fraction of the incumbent's brand effect.²⁸ This effect is plugged into the indirect utility function of the incumbent's product under the counterfactual policy (V_{ij}^{CTF}), and consumer surplus is calculated as we describe above.

The estimated upper bound of annual changes in consumer surplus are reported in Table 5. We find that for a relatively modest impact of the information treatment—inducing 25 percent of the customers to search and changing the perception of the new retailer brand to be 25 percent of the brand effect of the incumbent—the per household consumer surplus would rise by \$50/year. The gains in consumer surplus rise in both the fraction searching and the new entrant brand effect, especially the brand effect. If 50 percent are induced to search and the new entrant brand effect is half of the incumbent brand effect, the per household gains are \$149/year. And if 75 percent are induced to search and the new entrant brand effect is three-quarters of the incumbent brand effect, the per household gains are \$309/year. We do not take a position on what fraction of households would read and respond to the bill insert. While survey evidence suggests that some electricity consumers are not

²⁷ Under this interpretation of the logit discrete choice model, the logit shock is interpreted as product characteristics observable to the consumer but not to the researcher.

²⁸ Rather than inflate the new entrant brand effect up to the incumbent brand effect, one could deflate the incumbent brand effect down to the entrant's effect, which would decrease the consumer surplus gain. In order to estimate an upper bound, we use the former approach.

TABLE 5—POLICY COUNTERFACTUALS: ANNUAL CHANGES IN CONSUMER SURPLUS PER HOUSEHOLD

	Entrant brand effect relative to incumbent				
	0 percent	25 percent	50 percent	75 percent	100 percent
Fraction searching					
2 percent	\$0 (status quo)	\$1	\$3	\$6	\$8
25 percent	\$32	\$50	\$73	\$101	\$133
50 percent	\$67	\$103	\$149	\$205	\$268
75 percent	\$101	\$155	\$225	\$309	\$404
100 percent	\$136	\$208	\$301	\$412	\$539

Notes: This table shows the average annual change in consumer surplus per household under the counterfactual experiment described in Section V. The status quo is that 2 percent of the incumbent’s customers search for alternative retailers in a given month and that new entrant retailers enjoy none of the incumbent’s brand effect. Under each counterfactual, we simultaneously change two characteristics of high inertia customers, i.e., those who are still purchasing from the incumbent after two years of retail choice. First, we increase the fraction of incumbent customers who search in a given month (displayed down the rows). Second, we endow the new entrant retailers with a fraction of the incumbent brand effect (displayed across the columns). The dollar figures are the estimated yearly increases in consumer surplus for each incumbent customer, calculated as discussed in Section V.

Source: Authors’ calculations

actively engaged in energy-related decision making, other research suggests that both repeated and one-time information interventions affect energy use behavior.²⁹

In our lower bound, we calculate consumer surplus changes where the incumbent brand effect is modeled as distorting choices but not affecting ex post utility. If the incumbent brand effect distorts choice, then consumers will choose the incumbent despite lower priced entrants being available. In this counterfactual, we eliminate any product differentiation (via both the observed brand effect and unobserved logit shocks). The rationale for eliminating any such product differentiation lies in the spirit of Bernheim and Rangel (2009) that standard choice-theoretic models used to generate positive descriptions of behavior are not always suitable for welfare evaluations when a feature of the choice environment affects decision making but may not be relevant for a social planner’s welfare evaluation.³⁰ Under the assumption that there are no product characteristics that impact utility, then retail electricity is a homogeneous product and the consumer benefit to switching retailers is merely the cost savings. In the lower bound calculation of the consumer surplus changes, consumers choose the lowest priced retailer whenever they search. If the information intervention induces 75 percent of incumbent customers to search, then the consumer surplus gain is \$96/year. If 50 percent of incumbent customers respond by searching, then the consumer surplus gain is \$64/year.

It is beyond the scope of this paper to model the change in pricing as a response to consumers who search more frequently or attribute less of a relative brand

²⁹There is ample evidence of barriers to engagement—Competition and Markets Authority (2015) surveys UK households and finds that 34 percent of customers report never considering switching retailers, and 36 percent do not think it is or know if it is possible to switch retailer, tariff, or payment method. However, Navigant (2013) finds that over half of households read and respond to OPOWER Home Energy Reports; Seattle City Light (2014) finds similar figures for reading newsletters. Moreover, both repeated interventions (Allcott and Rogers 2014) and one-time information interventions (Pellerano et al. 2017) have been shown to affect energy consumption.

³⁰By the logic of Bernheim and Rangel (2009), analysts may view an institutional feature of the market—such as consumer misunderstanding of the mechanics of power flow under retail choice—as an ancillary condition of the choice environment rather than a welfare-relevant characteristic of the product.

advantage to the incumbent. However, it is worthwhile to keep in mind that the incumbent's price-to-beat is regulated and cannot be changed in response to any change in demand-side behavior. We leave an equilibrium model of the entrant retailers' supply-side response to future work.

This counterfactual experiment suggests that a relatively low-cost policy intervention—adding a flyer to an existing bill—could meaningfully increase consumer surplus. We should emphasize that we are modeling only changes in consumer surplus—each household is continuing to purchase the same amount of power each month, so our information treatment leads to a reallocation of rents from the incumbent to the consumers and new retail firms. Nevertheless, these results suggest that low-cost interventions to reduce choice frictions can meaningfully increase the consumer benefits of retail choice policy.

VI. Conclusions

This paper investigates an important policy change that brings choice to consumers in a previously regulated setting. We find that even in a relatively homogeneous good market, factors other than price competition play a very important role in shaping market outcomes: after four years of deregulation, the incumbent still maintains over 60 percent market share, despite the fact that some competitors consistently offer lower prices. Our model shows that two sources of inertia are prevalent—households do not frequently consider offerings of alternative retailers and they attach a significant brand advantage to the incumbent, especially in the early years of retail choice. Moreover, households in neighborhoods with lower income, less education, and more senior citizens face more inertia in both dimensions.

This paper suggests that there may be low-cost information interventions that reduce both sources of inertia. These types of interventions are likely to be particularly valuable as policy instruments in the early years of retail choice in the many jurisdictions that are expanding choice. And given the evidence of distributional consequences of retail choice, information interventions could be targeted to specific neighborhood demographics that face larger choice frictions.

Residential electricity is just one setting where policymakers are shifting to a regime where customers are first provided with choice. Other types of electricity customers—commercial and industrial—also are being offered retail choice, as are all types of customers of natural gas. Even more broadly, households are increasingly provided with choice in health care, retirement, and education. The types of inertia that we study in the residential electricity market are likely to present, perhaps to differing degrees, in many of these settings. Our paper provides a model to quantitatively assess the magnitude of different mechanisms that drive inertia. In addition, our framework allows one to measure consumer surplus gains from policies that reduce the sources of inertia.

REFERENCES

- Abaluck, Jason, and Jonathan Gruber. 2011. "Choice Inconsistencies among the Elderly: Evidence from Plan Choice in the Medicare Part D Program." *American Economic Review* 101 (4): 1180–1210.
- Allcott, Hunt, and Todd Rogers. 2014. "The Short-Run and Long-Run Effects of Behavioral Interventions: Experimental Evidence from Energy Conservation." *American Economic Review* 104 (10): 3003–37.
- Anderson, Soren T., Ryan Kellogg, and James M. Sallee. 2013. "What do consumers believe about future gasoline prices?" *Journal of Environmental Economics and Management* 66 (3): 383–403.
- Bernheim, B. Douglas, and Antonio Rangel. 2009. "Beyond Revealed Preference: Choice-Theoretic Foundations for Behavioral Welfare Economics." *Quarterly Journal of Economics* 124 (1): 51–104.
- Borenstein, Severin, James B. Bushnell, and Frank A. Wolak. 2002. "Measuring market inefficiencies in California's restructured wholesale electricity market." *American Economic Review* 92 (5): 1376–1405.
- Brennan, Timothy J. 2007. "Consumer preference not to choose: Methodological and policy implications." *Energy Policy* 35 (3): 1616–27.
- Bushnell, James B., Erin T. Mansur, and Celeste Saravia. 2008. "Vertical Arrangements, Market Structure, and Competition: An Analysis of Restructured U.S. Electricity Markets." *American Economic Review* 98 (1): 237–66.
- Competition and Markets Authority. 2015. *Energy Market Investigation: Provisional Findings Report*. Department for Business, Innovation and Skills. London, July.
- Dafny, Leemore, Kate Ho, and Mauricio Varela. 2013. "Let Them Have Choice: Gains from Shifting Away from Employer-Sponsored Health Insurance and Toward an Individual Exchange." *American Economic Journal: Economic Policy* 5 (1): 32–58.
- De los Santos, Babur, Ali Hortaçsu, and Matthijs R. Wildenbeest. 2012. "Testing Models of Consumer Search Using Data on Web Browsing and Purchasing Behavior." *American Economic Review* 102 (6): 2955–80.
- De los Santos, Babur, Ali Hortaçsu, and Matthijs R. Wildenbeest. 2013. "Search with Learning." Unpublished.
- Duarte, Fabian, and Justine S. Hastings. 2012. "Fettered Consumers and Sophisticated Firms: Evidence from Mexico's Privatized Social Security Market." National Bureau of Economic Research (NBER) Working Paper 18582.
- Dube, Jean-Pierre, Gunter J. Hitsch, and Peter E. Rossi. 2010. "State dependence and alternative explanations for consumer inertia." *RAND Journal of Economics* 41 (3): 417–45.
- Erdem, Tülin, and Michael P. Keane. 1996. "Decision-Making Under Uncertainty: Capturing Dynamic Brand Choice Processes in Turbulent Consumer Goods Markets." *Marketing Science* 15 (1): 1–20.
- Giulietti, Monica, Catherine Waddams Price, and Michael Waterson. 2005. "Consumer Choice and Competition Policy: A Study of UK Energy Markets." *Economic Journal* 115 (506): 949–68.
- Giulietti, Monica, Michael Waterson, and Matthijs Wildenbeest. 2014. "Estimation of Search Frictions in the British Electricity Market." *Journal of Industrial Economics* 62 (4): 555–90.
- Handel, Benjamin R. 2013. "Adverse Selection and Inertia in Health Insurance Markets: When Nudging Hurts." *American Economic Review* 103 (7): 2643–82.
- Hastings, Justine S., Ali Hortaçsu, and Chad Syverson. 2013. "Advertising and Competition in Privatized Social Security: The Case of Mexico." http://home.uchicago.edu/~hortacsu/HHS_UCLA_022813.pdf.
- Hastings, Justine S., Thomas J. Kane, and Douglas O. Staiger. 2010. "Heterogeneous Preferences and the Efficacy of Public School Choice." people.virginia.edu/~sns5r/microwkshp/hastings.pdf.
- Hastings, Justine S., and Jeffrey M. Weinstein. 2008. "Information, school choice, and academic achievement: Evidence from two experiments." *Quarterly Journal of Economics* 123 (4): 1373–1414.
- Hong, Han, and Matthew Shum. 2006. "Using price distributions to estimate search costs." *RAND Journal of Economics* 37 (2): 257–75.
- Honka, Elisabeth. 2014. "Quantifying search and switching costs in the US auto industry." *RAND Journal of Economics* 45 (4): 847–84.
- Honka, Elisabeth, Ali Hortaçsu, and Maria Ana Vitorino. 2014. "Advertising, Consumer Awareness and Choice: Evidence from the U.S. Banking Industry." Unpublished.
- Hortaçsu, Ali, Seyed Madanizadeh, and Steven L. Puller. 2017. "Power to Choose? An Analysis of Consumer Inertia in the Residential Electricity Market: Dataset." *American Economic Journal: Economic Policy*. <https://doi.org/10.1257/pol.20150235>.

- Hortaçsu, Ali, and Steven L. Puller.** 2008. "Understanding strategic bidding in multi-unit auctions: A case study of the Texas electricity spot market." *RAND Journal of Economics* 39 (1): 86–114.
- Hortaçsu, Ali, and Chad Syverson.** 2004. "Product Differentiation, Search Costs, and the Welfare Effects of Entry: A Case Study of S&P 500 Index Funds." *Quarterly Journal of Economics* 119 (2): 403–56.
- Ito, Koichiro.** 2014. "Do Consumers Respond to Marginal or Average Price? Evidence from Nonlinear Electricity Pricing." *American Economic Review* 104 (2): 537–63.
- Joskow, Paul L.** 2000. "Why Do We Need Electricity Retailers? Or Can You Get It Cheaper Wholesale?" Center for Energy and Environmental Policy Research (CEEPR) Working Paper 00-001.
- Keane, Michael P.** 1997. "Modeling Heterogeneity and State Dependence in Consumer Choice Behavior." *Journal of Business and Economic Statistics* 15 (3): 310–27.
- Ketcham, Jonathan D., Claudio Lucarelli, Eugenio J. Miravete, and M. Christopher Roebuck.** 2012. "Sinking, Swimming, or Learning to Swim in Medicare Part D." *American Economic Review* 102 (6): 2639–73.
- Kim, Jun B., Paulo Albuquerque, and Bart J. Bronnenberg.** 2011. "Mapping Online Consumer Search." *Journal of Marketing Research* 48 (1): 13–27.
- Kling, Jeffrey R., Sendhil Mullainathan, Eldar Shafir, Lee C. Vermeulen, and Marian V. Wrobel.** 2012. "Comparison Friction: Experimental Evidence from Medicare Drug Plans." *Quarterly Journal of Economics* 127 (1): 199–235.
- Koulayev, Sergei.** 2014. "Search for differentiated products: Identification and estimation." *RAND Journal of Economics* 45 (3): 553–75.
- Luco, Fernando.** 2014. "Distinguishing Sources of Inertia in a Defined-Contribution Pension System." Unpublished.
- Miravete, Eugenio J.** 2003. "Choosing the Wrong Calling Plan? Ignorance and Learning." *American Economic Review* 93 (1): 297–310.
- Miravete, Eugenio J., and Ignacio Palacios-Huerta.** 2014. "Consumer Inertia, Choice Dependence and Learning from Experience in a Repeated Decision Problem." *Review of Economics and Statistics* 96 (3): 524–37.
- Moraga-González, José Luis, and Matthijs R. Wildenbeest.** 2008. "Maximum likelihood estimation of search costs." *European Economic Review* 52 (5): 820–48.
- Navigant.** 2013. *Home Energy Reports Program: Program Year 2012 Evaluation Report, Appendix H.* AEP Ohio. Chicago, May.
- Neilson, Christopher.** 2013. "Targeted Vouchers, Competition among Schools, and the Academic Achievement of Poor Students." https://economics.sas.upenn.edu/system/files/event_papers/Neilson_2013_JMP_current.pdf.
- Pellerano, José A., Michael K. Price, Steven L. Puller, and Gonzalo Sánchez.** 2017. "Do Extrinsic Incentives Undermine Social Norms? Evidence from a Field Experiment in Energy Conservation." *Environmental and Resource Economics* 67 (3): 413–28.
- Seattle City Light.** 2014. "Seattle City Light Community Energy Forum: Communications study." Unpublished.
- Shin, Sangwoo, Sanjog Misra, and Dan Horsky.** 2012. "Disentangling Preferences and Learning in Brand Choice Models." *Marketing Science* 31 (1): 115–37.
- Small, Kenneth A., and Harvey A. Rosen.** 1981. "Applied Welfare Economics with Discrete Choice Models." *Econometrica* 49 (1): 105–30.
- Sweeting, Andrew.** 2007. "Market Power in the England and Wales Wholesale Electricity Market 1995–2000." *Economic Journal* 117 (520): 654–85.
- Waddams Price, Catherine, Catherine Webster, and Minyan Zhu.** 2013. "Searching and Switching: Empirical estimates of consumer behaviour in regulated markets." Centre for Competition Policy (CCP) Working Paper 13-11.
- Wilson, Chris M., and Catherine Waddams Price.** 2010. "Do consumers switch to the best supplier?" *Oxford Economic Papers* 62 (4): 647–68.
- Wolfram, Catherine D.** 1999. "Measuring Duopoly Power in the British Electricity Spot Market." *American Economic Review* 89 (4): 805–26.